

d) What are the possible sequence of execution of the last printf() and also specify the values of n and m printed?

Ans: Two possible sequence of execution : \rightarrow



- ① If parent execute before C₂ : \rightarrow
- | | | | |
|--------------|----|----|--------------------------|
| last printf | 46 | 67 | (due to P) |
| last print f | 41 | 62 | (due to C ₂) |
- ② If child execute first (before parent) : \rightarrow
- | | | | |
|--------------|----|----|--------------------------|
| last printf | 41 | 62 | (due to C ₂) |
| last print f | 46 | 67 | (due to P) |

4. What happens when a child completes its execution before the parent executes a wait() system call? What happens when parent dies before the child finishes its execution.(assume Linux /UNIX OS)

Ans: ① In this case the execution will ~~not affect~~⁽²⁾ the function.
In any way. As wait system call ensures that ~~at least~~⁽²⁾ one of its parent starts executing before its parent exits.

② The child becomes orphan. And in Linux / UNIX OS ; the parent dies in this case ; the child becomes orphan. And its parent . So child orphans child processes have init process go to its parent ; returns its process id to init process .

5. Can a multithreaded solution using multiple user level threads achieve better performance on a multiprocessor system than on a single processor system?(2)

Ans: On a single processor system , execution will be sequential ; but on a multiprocessor system ; OS have certain options for execution like should all the processes / certain processes may run parallel . Having the operation (any specified operation) should achieve respond . For multi processor ; if it is efficient and has better performance .

6. Given the following job mix

Process	Arrival time	Burst time (in msec.)	Priority
P1	0	15	6
P2	0	10	4
P3	3	3	0 (Highest)
P4	7	2	2

Calculate the average turn around time and average waiting time for:

$$WT = \text{Waiting time}$$

a) A non-preemptive priority algorithms

Ans:

	P_2	P_3	P_4	P_1
0	10	13	15	30

$$\text{avg waiting time} = \frac{\text{waiting}(WT)P_1 + (WT)P_2 + (WT)P_3 + (WT)P_4}{4}$$

$$= \frac{(15-0) + (0-0) + (10-3) + (13-7)}{4}$$

$$= \frac{28}{4} = 7 \text{ msec}$$

avg turnaround time = CPU time + I/O time + waiting time (due to I/O device queue)

+ waiting time (due to ready queue)

~~$$= \frac{(15+15)+(10+0)+(3+7)+(2+6)}{4}$$~~
~~$$= \frac{58}{4} = 14.5 \text{ msec}$$~~

b) A preemptive priority algorithms.

Ans:

	P_2	P_3	P_4	P_1
0	3	6	7	9

$$\text{avg waiting time} = \frac{(15-0) + (9-3-1-0) + (3-3) + (7-7)}{4}$$

$$= \frac{15+5}{4} = \frac{20}{4} = 5 \text{ msec}$$

$$\text{avg turn around time} = \text{avg waiting time} + \text{avg CPU burst time}$$

$$= 5 + \frac{20}{4} = 5 + 7.5$$

$$= 12.5 \text{ msec}$$

actually turn around time = CPU time + I/O time + (burst) due to I/O device queue + (waiting time due to ready queue).

No info about I/O time

Roll No: 3130055CS

Name..... ASHOK DHONI GANJA



National Institute of Technology Calicut

Department of Computer Science & Engineering

CS3003 – OPERATING SYSTEMS

First MID Term Exam (Monsoon Semester 2015)

PART B

Max. Marks: 7

Time: 30 mnts

1. Which type of fragmentation(s) does the following memory management schemes suffer from?

Paging : Ans: ~~External fragmentation~~

Segmentation : Ans: ~~Internal fragmentation~~

2. Consider six memory partitions of size 200 KB, 400 KB, 600 KB, 500 KB, 300 KB, and 250 KB, where KB refers to kilobyte. These partitions need to be allotted to four processes of sizes 357 KB, 210 KB, 468 KB and 491 KB in that order. If the best fit algorithm is used, which partitions are NOT allotted to any process?

Ans: ~~200 KB, 300 KB~~

3. Given an effective access time of 300ns, a main memory access time of 250ns, a TLB access time of 20ns what is the hit rate? $\text{Hit rate} = \frac{(1-x)}{(1-x) + (20 + 250 + 250)}$

Ans: $300 = (20 + 250)x + (1-x)(20 + 250 + 250)$

$$300 = 270x + (1-x)520$$

$$250 = 220 \quad ; \quad x = \frac{220}{250} = 0.88 \quad ; \quad \text{Hit rate} = 0.88$$

4. Consider the virtual page reference string 7,1,2,4,5,1,2,4,3,7,5,8,3,2,7,3,8,2,1,1,7,4,7,2,8,4,7,1,5. On a demand paged virtual memory system running on a computer system that main memory size of 4 pages frames which are initially empty. Let LRU, FIFO and OPTIMAL denote the number of page faults under the corresponding page replacements policy. Then choose the right relation and How it comes

a. OPTIMAL < LRU < FIFO

b. OPTIMAL < FIFO < LRU ✓

c. OPTIMAL = LRU

d. OPTIMAL = FIFO

(b)

$$\begin{array}{l}
 \text{FIFO} \quad PF = 14 \\
 \text{LRU} \quad PF = 15 \\
 \text{OPTIM} \quad PF = 11
 \end{array}$$

Ans:

2

Given a 32 MB memory and a frame size of 2048 bytes.

二十一

following address in terms of page number and offset within the page:

[2]

Ans: 1 2 0

$$\text{No. of frames} = \frac{32 \times 2}{2} = 16 \times 2^4 = 2^4 \times 2^4 = 2^{11}$$

210

$$2^{32} / 18^{10} = 2^{21} \approx$$

bikes are used to ride into page table

三

21 bits are used to index into page table
1 bit is used for page offset.

11

Given address

卷之三

11 bits
11101010000000001001

$$\begin{aligned} \text{page number} &= \cancel{45} - 25 = \\ \text{page offset} &= 11 \end{aligned}$$

Name BHOPALKA

Roll No: .B130057CS.....



National Institute of Technology Calicut

Department of Computer Science & Engineering

CS3003 – OPERATING SYSTEMS

First MID Term Exam (Monsoon Semester 2015)

PART A

1. Assume you have following jobs to execute with one processor. Suppose system uses RR

Scheduling with time slice as 10 units. Assume context switch time is 3 units.

Pi	Service time(Pi)
0	85
1	30
2	20
3	35
4	45

a) What is the turnaround time for the process P3?

Ans: Turnaround time = 200 units

0	13	26	39	52	65	78	91	104	117	130	143	156	169	182	195	208
P0	P1	P2	P3	P4	P0	P1	P2	P3	P4	P0	P1	P2	P3	P4	P0	P1

b) What is the average wait time P3?

Ans: 39 units

[1 mark]

c) What is the effect of increasing the time quantum to an arbitrarily large number for Round Robin scheduling?

[0.5 mark]

Ans: The turnaround time and the wait time for the processes will increase

Fair-share scheduler

2. Linux 2.6.23 kernel usesScheduler which uses~~doubly linked list~~.....~~data structure~~..... [1.5 marks]

3. System calls are function calls made from user space programs into the kernel requesting some services. How will the system pass system call number from user space to kernel space?

Ans: With the help of registers

4. Fork returns~~parent process id~~..... to parent process on success [0.5 mark]

Ans:

5. How many times printf() will be executed in the below mentioned program?

```
main() {  
    int i;
```

```
    for (i = 0; i < 4; i++)
```

```
    fork();
```

```
    printf("pid = %d\n", getpid());  
}
```

[1 Mark]

Ans: 6 times

6. What are the process states in Linux

Ans: ready to run
~~sleep~~

7. A~~process~~..... is a process that is still executing, but whose parent has died.

Ans:

8. A process which has finished its execution, but its parent has not yet called to collect its return status, becomes a~~combine~~..... [0.5 Mark]

Ans:

Max. Marks: 7
Time: 30 mnts

1. Which type of fragmentation does the following memory management schemes suffer from?

Paging : Ans: ~~External~~

Segmentation : Ans: ~~internal & external~~

2. Consider six memory partitions of size 200 KB, 400 KB, 600 KB, 500 KB, 300 KB, and 250 KB, where KB refers to kilobyte. These partitions need to be allotted to four processes of sizes 357 KB, 210 KB, 468 KB and 491 KB in that order. If the best fit algorithm is used, which partitions are NOT allotted to any process?

Ans: ~~Ans: 250 KB → 400 KB
210 KB → 250 KB
468 KB → 500 KB
491 KB → 600 KB~~

~~So, 200 KB, 300 KB
won't be allotted to any process.~~

3. Given an effective access time of 300ns, a main memory access time of 250ns, a TLB access time of 20ns what is the hit rate?

Ans: Effective access time = $(hit\ ratio) * TLB\ access\ time + (1 - hit\ ratio) * (Main\ Memory\ access\ time + TLB\ access\ time)$

$$300 = \alpha * 20 + (1 - \alpha)(250) \Rightarrow 300 = 20\alpha + 250 - 250\alpha$$

$$\Rightarrow 300 = -250\alpha + 250 \Rightarrow \alpha = \frac{250}{250} = 1$$

4. Consider the virtual page reference string 7,1,2,4,5,1,2,4,3,7,5,8,3,2,7,3,8,2,1,1,7,4,7,2,8,4,7,1,5. On a demand paged virtual memory system running on a computer system that main memory size of 4 pages frames which are initially empty. Let LRU, FIFO and OPTIMAL denote the number of page faults under the corresponding page replacement policy. Then choose the right relation and How it comes

a. OPTIMAL < LRU < FIFO

b. OPTIMAL < FIFO < LRU

c. OPTIMAL = LRU

d. OPTIMAL = FIFO

Strategy for optimal ~~best~~ fastest in future,

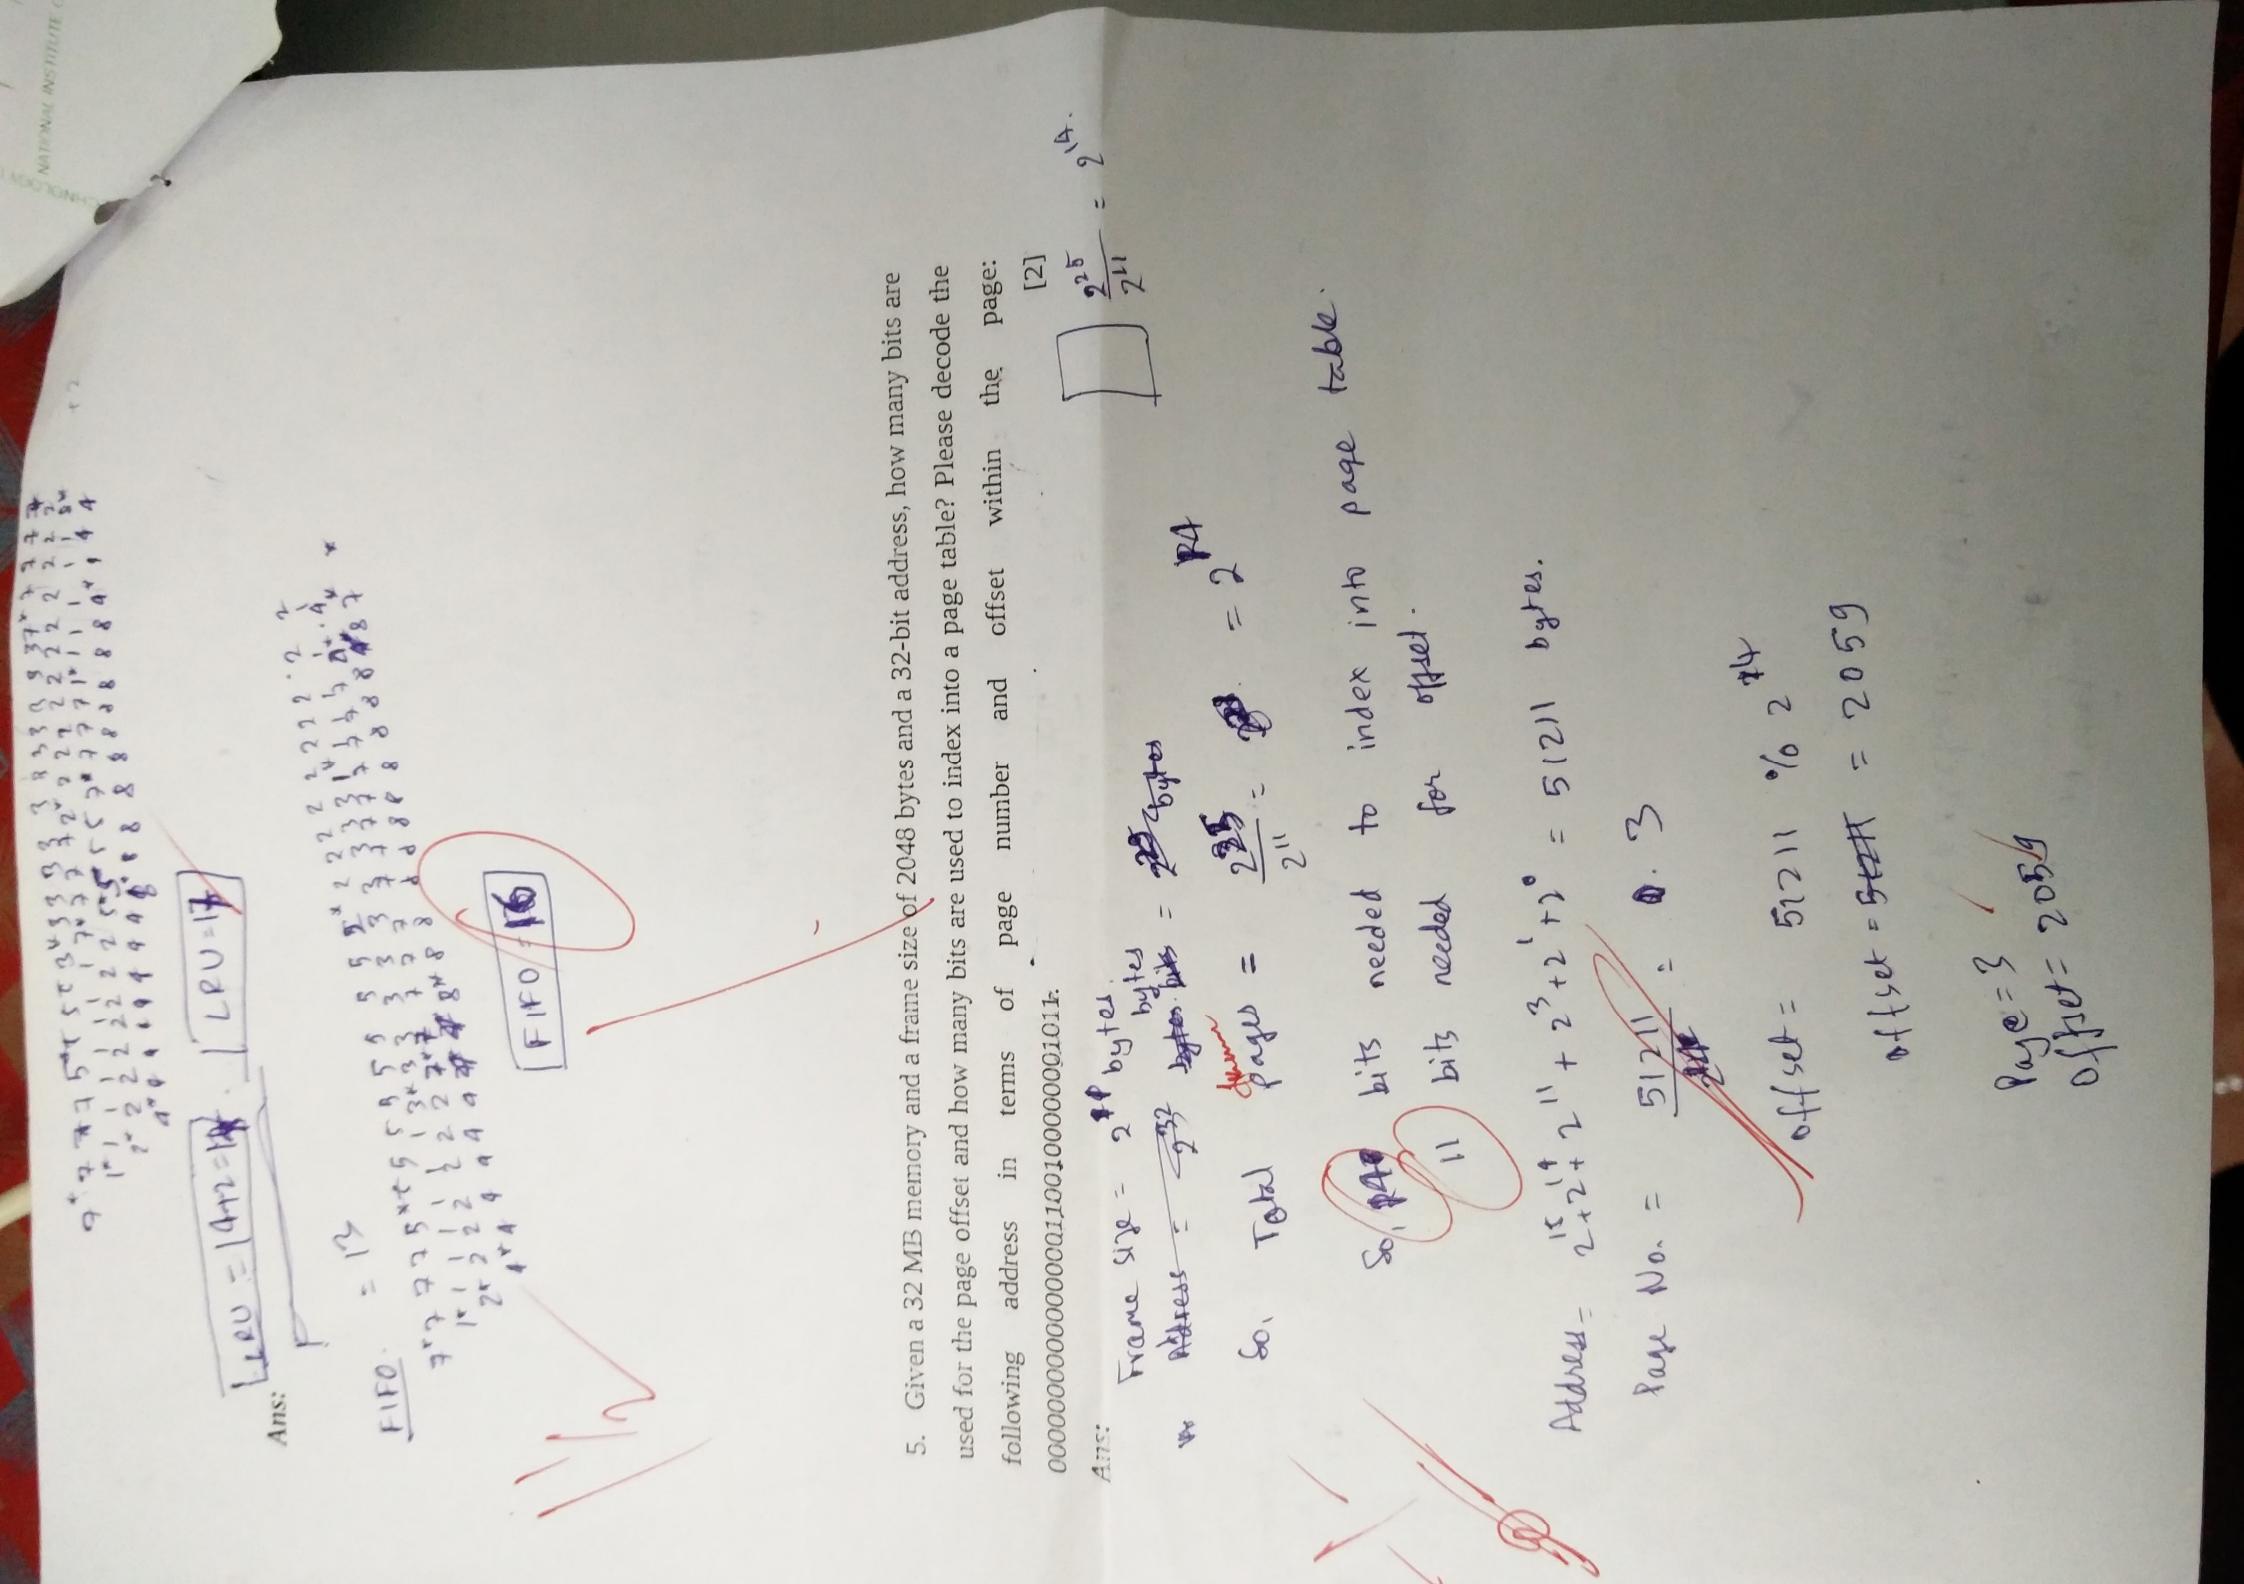
OPTIMAL

LRU

7 1 2 4 5 1 2 4 3 3 5 8 3 2 3 3 2 1 3 4 3 2 2 4 7 1 5

OPTIMAL = 12

A + 2 + 1 + 1 + 1 + 1 + 1 + 2





Department of Computer Science & Engineering
NATIONAL INSTITUTE OF TECHNOLOGY CALICUT
Monsoon Semester 2015
End Semester Examination
CS 3003 OPERATING SYSTEMS

Roll No.....

Duration: 1hour30mins

PART B

Maximum Marks: [25]

(1) Name the following:

- a. Frequent page faults results in an undesirable condition such as slowing of memory [1]
- b. The field in *Segment Selector* specifies whether the *Segment Descriptor* is included in the *GDT* or in the *LDT* [1]

2. Define following terminologies

- a. reentrant code
 - b. double buffering
 - c. inverted hash table
3. Draw the block diagram of cascade mode of 8259 controller [1]
4. Write the RAID level corresponding to each of the following:
- a) Block level distributed parity : RAID __
 - b) Mirrored: RAID __
 - c) Non redundant: RAID __
 - d) Hamming Code: RAID __

5. Consider a system with 2 level paging scheme in which a regular memory access takes 200ns and servicing a page fault takes 10ms. An average instruction takes 100ns of CPU time and 2 memory access. The TLB hit ratio is 95% and the page fault rate is one in every 10,000 instructions. What is the effective average instruction execution time? [2]
- (6) Consider I/O based polling and memory mapped I/O. In what situation would you favour one technique over the other? [1.5]

7. Consider four frames with following page references
A,B,A,C,D,E,F,A,B,C,F,A,E,F,B,D,E,C,A,F
Show the content of the frames at every access for the following page replacement policies and Find out the number of page faults. [2]

- a. Random & optimal algorithm b. LRU
8. Suppose you have a 4²-bit virtual address space with a page size of 16 KB and that page table entry takes 8 bytes. How many levels of page tables would be required to map the virtual address space if every page table is required to fit into a single page? Be explicit in your explanation and show how a virtual address is structured.

9. The following table represents three different possible designs for memory.

Characteristic	Design 1	Design 2	Design 3
Physical memory address width	8bit	16bit	32bit
Logical address width	12bit	20bit	24bit
Page size	16bytes	32bytes	64bytes

- a) For each design, find the maximum number of pages each process can have in logical address space.
- [1]
- b) Which design is likely to have the highest internal fragmentation (if any)?
- [1]
- c) Which design is likely to have the highest external fragmentation (if any)?
- [1]
10. Briefly describe the steps taken to read a block of data from the disk to the memory using DMA controlled I/O.
- [2]
11. Provide the detailed steps involved in performing an I/O operation using interrupt mechanism in an 8086 architecture
- [2]
12. Suppose that we have a 64-bit virtual address uses segmented paging, which split as follows:
- | | | | | |
|------------------------------|-----------|-----------|-----------|----------|
| 6bits | 11bits | 11bits | 11bits | 14bits |
| [Segment 16bit]
[Level 1] | [Level 2] | [Level 3] | [Level 4] | [Offset] |
- a. How many segments are in this system?
- [0.5]
- b. Assume that the page tables are divided into page-sized chunks (so that they can be paged to disk). How much space have we allowed for a PTE in this system?
- [0.75]
- c. What is the page table overhead for a process that uses one segment? Explain.
- [1.25]



Department of Computer Science & Engineering
NATIONAL INSTITUTE OF TECHNOLOGY CALICUT
Monsoon Semester 2015
End Semester Examination
CS 3003 OPERATING SYSTEMS

Duration: 1hour30mmts

PART A

- Maximum Marks: [25]**
1. Suppose a system uses priority scheduling (under the following process load), where a small integer means a high priority

[2.5]

Pi	Service time(Pi)	Priority
P0	80	3
P1	20	1
P2	10	4
P3	20	5
P4	50	2

Compute the following

- Create a Gantt chart illustrating the execution of these processes
- What is the turnaround time for process P2 under priority scheduling?
- What is the average wait time for the processes?

2. How many numbers of processes will be spawned after executing this program? Justify your answer

#include <stdio.h>

```
#include <unistd.h>
int main()
{
    fork();
    fork() && fork() || fork();
    fork();
    printf("forked\n");
    return 0;
}
```

3. What is context switching? Suggest an idea to reduce the context switching time?
4. Draw the process state diagram
5. Suppose we want to synchronize two concurrent processes P and Q using binary semaphores S and T
- The code for process P and Q is shown below: Synchronization statements can be inserted only at points W,X,Y,Z. Which of the following will always lead to an output starting with 001100110011?

PROCESS P	PROCESS Q
While(1){ W: print Q; print Q; X: }	While(1){ Y: print I; print I; Z: }

- a) P(S) at W,V(S) at X,P(T) at Y,V(T) at Z, S=T=1 initially
 b) P(S) at W,V(T) at X,P(T) at Y,V(S) at Z, S=1 and T=0 initially
 c) P(S) at W,V(T) at X,P(T) at Y,V(S) at Z, S=0 and T=1 initially
 d) P(S) at W,V(S) at X,P(T) at Y,V(T) at Z, S=T=0 initially
- [1]
6. Compare a counting semaphore and a binary semaphore
7. A semaphore s is a nonnegative integer variable changed or tested only by one of two indivisible access routines V(s) & P(s). Define V(s) and P(s) functions for a data type semaphore S;
- [2]
8. The test-and-set instruction (TS) is an instruction used to write to a memory location and return its old value as a single atomic. Implement a general semaphore using this test and set instruction! (P & V functions of semaphore)
- [2]
9. What are the necessary conditions for a deadlock to happen?
10. Consider a disk in which blocks 2, 3, 4, 6, 8, 9, 10, 11, 12, 15, 17, 19, 25, 26, and 28 are free and the rest of the blocks are allocated. Prepare a bit vector (for free space management) for the above scenario?
- [1]
11. Answer the following
- a. What is i-node(index node) in Unix
 b. What will be content of a i-node, corresponding to a device? Brief them
12. Suppose the disk queue contains a set of references for blocks on tracks 76,124,17,269,201,29,137 and 12. Assume disk head is currently pointing to cylinder 57. Perform the SCAN Disk optimization and find how many tracks the head will move across as it serves the request
13. Banker's algorithm is the best known of the deadlock avoidance strategies. Given following details such as alloc table, Resource type C, maximum claim table. Find given allocation is safe or unsafe with the help of banker algorithm. If safe provide the order of execution also
- [3]
- | MaxC | Allocation | Available Resource Type – C |
|------------|------------|-----------------------------|
| ABCD | ABCD | |
| P0 3 2 1 1 | 2 0 1 0 | A B C D |
| P1 1 2 0 2 | 1 1 0 0 | 6 4 4 2 |
| P2 1 1 2 0 | 1 1 0 0 | |
| P3 3 2 1 0 | 1 0 1 0 | |
| P4 2 1 0 1 | 0 1 0 1 | |
- [0.5+0.5]
14. What is the importance of VFS switch in Linux?
15. Mention 3 different file allocation techniques
16. Explain the algorithm used for Open() system call with the help of the necessary datastructure [1.5]
17. Draw the sequence of client server establishment using socket related system call
- [1]

* For Questions 1-3 identify if the statements given are TRUE or FALSE. Marks will be awarded only if proper justification is given.

1. In a demand paging system, if a virtual page number P generates a TLB miss, then the corresponding frame number for P is guaranteed to be found in the Page Table Entry. (1 mark)

~~TRUE → Page Table of process contains the address of page fault when we have to bring a page from disk to main memory & then update the entry in PT.~~ (1 mark)

2. The accuracy of working set depends on the window size. (1 mark)

~~$WS = \min(D, N) \rightarrow \text{no. of elements (distinct) elements in a set.}$~~

3. Adding pre-fetching to page-fetch algorithms can increase CPU utilization. (1 mark)
~~FALSE → Pre-fetching reduces the page fault rate for a particular process. To increase CPU utilization, we must increase the degree of multi-programming.~~

4. Considering a page size of 1 KB and that each page table entry takes 4 bytes, how many levels of page tables will be required to map a 34-bit address if every page table fits into a single page. Clearly justify your answer. (2 marks)

$$\begin{aligned} \text{Page size} &= 1 \text{ KB} \\ \text{no. of bytes (inner)} &= 2^{34} / 1 \text{ KB} = 2^{24} = 16M \\ \text{Page table size} &= 16 \times 4 = 64 \text{ M} \rightarrow \text{more than size of page table} \\ \text{Top-level}(\alpha) &= \frac{64M}{128} = 2 \cdot 2^{20} = 64K \rightarrow \text{more (entries)} \\ \text{Top level(b)} &= 256 \text{ K} / 1K = 256 \text{ bytes (entries)} \\ \text{size} &\approx 256 \times 4 = 1024 \text{ bytes} = 112 \rightarrow \text{bytes} \end{aligned}$$

~~3 levels~~ ✓ → is the total size of page is in memory.

5. In a demand paged memory, 200 nanoseconds are required to satisfy a memory request if page is in memory. It takes 8 milliseconds to service a page fault if an empty frame is available or if the page to be replaced is not modified; and 20 milliseconds if the page to be replaced is modified. 60% of the time a page fault occurs, the page to be replaced is dirty. What is the effective access time if the page fault rate is 10%?

$$\begin{aligned} T_{mem} &= 200 \text{ ns} \\ T_{pf} &= 8 \text{ ms} \\ T_{effective} &= [1 - 0.1] [2 \cdot T_{mem}] + 0.1 [60\% (T_{pf} + 2 \cdot T_{mem}) + 40\% (T_{pf} + 2 \cdot T_{mem})] \\ &= (1 - 0.1) [2 \cdot 200 \text{ ns}] + 0.1 (0.6 (20 \text{ ms} + 2 \cdot 20 \text{ ms}) + 0.4 (8 \text{ ms} + 2 \cdot 20 \text{ ms})) \\ &= 368.84 \text{ ms} \\ &= 368.84 \times 10^{-3} \text{ s} \\ &= 0.36884 \text{ s} \end{aligned}$$

~~= 36000 + 0.1(200024 + 3,20016 \times 10^{-3})~~
~~= 0.012320616 \times 10^{-3}~~

$$= 1208$$

$$T_S = 0.9 \times 3 \text{ bytes} + 0.1 [0.01200024 + 3,20016 \times 10^{-3}] \\ = 3.6 \times 10^{-3} + 0.1 (0.01200024 + 3,20016 \times 10^{-3}) \\ = 1.5204 \times 10^{-3}$$

6. Consider the following piece of code whose function is to initialize each element of a 512-by-512 integer array A to 0. (The array is stored in Row-major order):

Initialize_1()

```
{ for (int i = 0; i < 512; i++)
    for (int j = 0; j < 512; j++)
        A[i][j] = 0;
```

Assume that the code for this function fits in one page, and the stack also fits in one page. Let the size of an integer be 4 bytes. Also, the page size is 4096 and the TLB has 4 entries. (Assume LRU replacement policy)

i) Calculate the number of TLB misses.

Code - 1 page
 stack data - 1 page
 data = ~~4x 512 = 2048 = 2 pages~~
 total = ~~512 x 512 x 4 = 1048576 bytes = 256 pages~~
~~entries of 2 rows~~ ~~number present at a time~~
~~Therefore, No. of TLB misses =~~
 1 new data = ~~512 x 4 = 2048~~
~~i. at a time total data 2 new data can be in 1~~
~~page. Since 2 pages or vacant for data.~~
 TLB misses = ~~512 / 4 = 128~~

ii) Calculate the number of TLB misses if the code is as given below, considering all the assumptions above.

Initialize_2()

```
{ for (int i = 0; i < 512; i++)
    for (int j = 0; j < 512; j++)
        A[i][j] = 0;
```

~~Y~~ Page fault occur more frequently no. Due to
~~A[0][i]~~ ~~this takes data from every other row and cause~~
~~page fault in every 4th execution~~

Total Misses =

$$(512 \times 512) / 4 \\ = 65536$$

~~National Institute of Technology Calicut~~

~~Department of Computer Science & Engineering~~

CS3003 – OPERATING SYSTEM

Second MID Term Exam (Monsoon Semester 2014)

Max. Marks: 10

PART A

1. a) Test and Set instruction reads the contents of the given memory word *lock* into the register RX and then stores a nonzero value at the memory address lock (TS RX, LOCK). How can this instruction be used to prevent two processes from simultaneously entering their critical regions. Write an assembly language subroutine (X86 based) to implement the P() function of

a semaphore used under a while condition [1 mark]

Test and set is used under a while condition. If 'm' was false initially like in Test and set (m) . If 'm' was allowed to execute then m is set to true and next msh is allowed to execute. Then other process will go in empty loop until meanwhile the other process can execute their operation and after that only they can execute their operation.

Assembly code

```
mov eax, [lock]
dec eax
mov S, [lock]
```

- b) An alternative instruction to TS is XCHG, which exchanges the contents of two locations atomically. Build an assembly language subroutine to implement the P() utilizing XCHG instruction (XCHG RX, LOCK)

```
mov eax, RX
mov ebx, LOCK
xchgl eax, RX
mov RX, [ebx]
mov LOCK, [eax]
```

2. Write 2 examples of socket communication applications

1. Distributed systems
2. Server - Client system

[1 mark]

3. Shared memory IPC shares the memory in user...space!

4. Specify two important advantage of thread over process?

- ① Process creation is costlier than thread creation.
- ② Context switch in case of threads is faster.

[0.5 mark]

[1 mark]

5. Write pseudo code to simulate & solve following problems using semaphore
- a) A bridge which allows traffic in both directions has a maximum capacity of 3 vehicles at a time. There are vehicles from Both sides(SideA, SideB) competing to pass through the bridge.

```

Semaphore max = 3;
SideA vehicles SideA {
    P(max)
    : //pass
    V(max)
}
SideB vehicles SideB {
    P(max)
    : //pass
    V(max)
}
    
```

- b) Perform data communication through a half duplex channel of capacity 5(packets). The channel is used by two processes. Both the processes(process1 and process2) are only sending the packets through the channel. Make sure that the following conditions are met during the data communication:

- 1) The process1 sends data in one direction and the process2 sends data in another direction of the channel.
- 2) If the channel is full or it is not the process's direction of communication, the process waits.

~~Half - duplex → Sending in one direction~~

```

#include < stdio.h>
#include < sys/types.h>
#include < stdio.h>
#include < stdio.h>

int main () {
    int id [2];
    int id1 = mknod ("buffer1", 0777, 0);
    id1 = Open ("process1", O_WRONLY);
    int id2 = mknod ("buffer2", 0777, 0);
    id2 = Open ("process2", O_RDONLY);
    buf packet [5];
    buf receive [5];
    int a = write (id1, packet, 5);
    close (id1);
    a = read (id2, &receive, 5);
    close (id2);
    a = close (id1);
    
```

6. Which signal will be sent to kill a process forcefully?

SIGKILL

[0.5 mark]

7. Write three ways to register and handle a signal in Linux

signal (signal - int , SIG - INT)
signal (signal - int , SIG - DEF)
signal (signal - int , Signal - handler)

[0.5 mark]

8. Write 2 different ways to delete a message queue whose id is 65536 [0.5 mark]

① msgctl (65536 , IPC - RMID , 0);
② mqdlt (65536);

[0.5 mark]

9. Write a C program to catch the hang up signal send by the terminal controller, and then print the message "IT IS A HIGH PRIORITY JOB" for once. Application should hang up from second time onwards when you close the terminal. [1.5 Marks]

```
#include < stdio . h >
#include < signal . h >
void * signal - hand ( int signumber ) {
    void * signal - hand ( int signumber ) {
        printf (" IT IS A HIGH PRIORITY JOB ");
        sleep ( 2 );
        signal ( SIGHUP , SIG - DEF );
        signal ( SIGHUP , SIG - DEF );
    }
}
int main ( int argc , int args [ 3 ] {
    int main ( int argc , int args [ 3 ] {
        signal ( SIGHUP , signal - handler );
        signal ( SIGHUP , signal - handler );
        return 0;
    }
    while ( 1 )
        sleep ( 1 );
}
```

sleep

National Institute of Technology Calicut

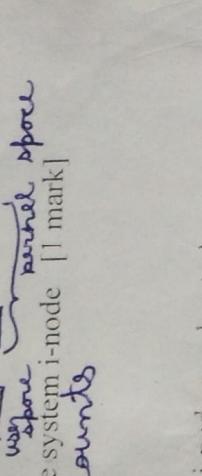
Department of Computer Science & Engineering

CS3003 - OPERATING SYSTEM

First MID Term Exam (Monsoon Semester 2014)

PART A

Time: 1 Hour

1. a) For user mode to kernel mode switching, Linux OS uses. calls [0.5 + 0.5 Mark]
b) How does the parameter from user area pass to kernel area while executing the system call using registers eax, ebx, ecx, edx, esi, edi
2. Open() system call return an integer number in case of a successful file opening.
a) What does this number indicate?
file descriptor index in file descriptor table [0.5 + 0.5 mark]
- b) How does kernel map this number to actual data block? using i-node table entry 
3. Which are the extra fields that an incore-iinode keep than a file system i-node [1 mark]
status - locked / unlocked | reference counts | i-node number
4. algorithm converts the file name to its corresponding i-node number to locate the data block in the file system. namei [0.5 mark]
5. Register context of a process is comprised of values in the registers namely SP, I, R, PC, etc. [1 mark]
6. Write a C program to implement a cp command using system calls

```
#include < stdio.h>
#include < unistd.h>
```

```
int main( ) {
```

```
char *from = "Inputfile"; // assumed to be source
char *to = "output"; // of copy
int id1 = open ("from", O_RDONLY);
int id2 = open ("to", O_WRONLY);
char buffer[5];
while (id1 != EOF) {
    read (id1, &c1, 1);
    write (id2, &c2, 1);
}
```

```
close (id1);
close (id2);
return 0;
```

7. Predict the output of the following

```
a) #include<stdio.h>
#include<stdlib.h>
#include<unistd.h>

int main()
{
    pid_t child;
    child = fork();
    switch(child){
        case -1:
            perror("fork");
            exit(1);
        case 0:
            sleep(10);
            printf("%d\n", getppid());
            break;
        default:
            break;
    }
    return 0;
}

b) #include<stdio.h>
#include<stdlib.h>
#include<unistd.h>

int main()
{
    pid_t child;
    int status;
    child = fork();
    switch(child){
        case -1:
            perror("fork");
            exit(1);
        case 0:
            printf("%d\n", getpid());
            break;
        default:
            printf("%d\n", getpid());
            wait(&status);
            break;
    }
    return 0;
}
```

*parent_id is printed.
also child becomes orphan here.*

op → parent_id

8. Assume you have following jobs to execute with one processor. Suppose system uses RR Scheduling with time slice as 10 units. Assume context switch time is 3 units.

Pi	Service time(Pi)
0	75
1	40
2	25
3	20
4	45

a) What is the turnaround time for the process P3?



[1 mark]

114 units

b) What is the ~~wait time~~ wait time for P3?

39 units

[1 Mark]

c) What is the effect of increasing the time quantum to an arbitrarily large number for Round Robin scheduling?

Both turn-around time and avg. wait time would increase.

[0.5 Mark]

NATIONAL INSTITUTE OF TECHNOLOGY CALICUT
 First Midterm Examination Monsoon Semester 2014
 CS3003 – OPERATING SYSTEMS

Max Marks : 9

PART - B

Answer in the space provided

1. Consider a system with 32-bit logical addresses, 16-bit physical addresses, and 1KB page size. Compute the Page Table size in bytes if each page table entry takes 4-bytes. Justify your answer (1 mark)
- ~~1KB = 2^{10} bytes, if each entry takes 4 bytes then, total no. of entries in each page = $2^{10} / 2^2 = 2^8$; B-bits are needed for offset we are left with $16 - 8 - 32 - 8$ bytes = 24 bit size in bytes = $2^{24} \times 1 \times 2^{10} = 2^{34}$ bytes~~

2. Consider a segmentation-based system. At some point in the system operation, the main memory has holes (ie, free segments) in this order : 20KB, 15KB, 90KB, 54KB, 10KB, 30KB, 56KB. A new request for memory of size 22KB arrives. Which hole will be allocated in each of the schemes? Justify your answer. (1 mark)

A) Best-Fit :

30 KB hole.

~~Reason :- Internal fragmentation is least in this case. Since it is just enough to behalf in 30KB hole.~~

B) First-Fit :

~~90 KB~~ It is the first block which can accommodate

22 KB

3. Which type of fragmentation(s) does the following memory management schemes suffer from? (2 marks)

A) Paging -

~~internal fragmentation~~

B) Segmentation -

~~internal fragmentation~~

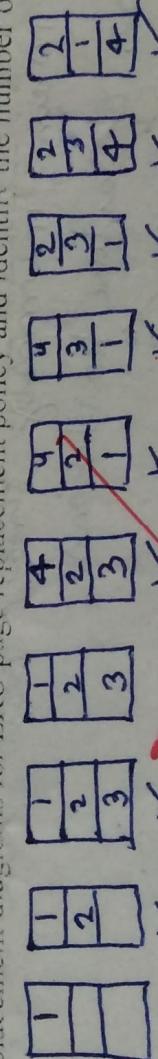
C) Buddy systems -

~~internal fragmentation as well as external.~~

D) Fixed Partitioning-

~~internal and external fragmentation.~~

4. Consider a virtual page reference string: 1, 2, 3, 2, 4, 1, 3, 2, 4, 1 for a demand paging system running on a computer with main memory size of 3 page frames, which are initially empty. Draw the frame allocation-replacement diagrams for LRU page replacement policy and identify the number of page faults. (2 marks)



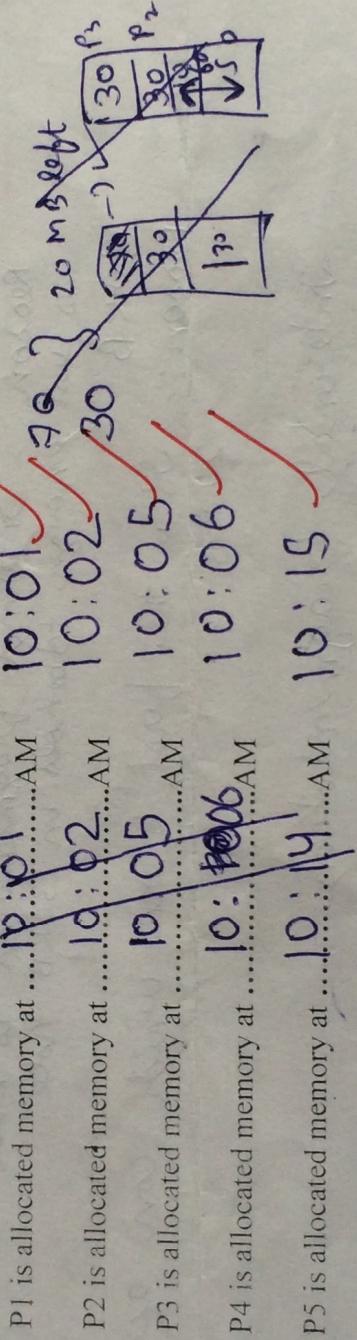
~~no. of page faults = 9.~~

5. Consider a system with 128MB of physical memory, of which the first 8MB is allocated to kernel. The remaining 120MB can be allocated to processes in contiguous chunks (dynamic partitioning). The creation times, finishing times, and memory requirements of 5 processes are given below:

PROCESS	START TIME	FINISH TIME	MEMORY (in MB)
P1	10:01AM	10:05AM	70
P2	10:02AM	10:10AM	30
P3	10:04AM	10:14AM	30
P4	10:06AM	10:15AM	10
P5	10:12AM	10:25AM	100

Assume the requests are granted in FIFO order according to First-fit algorithm and each process occupies memory allocated to it until it terminates.

- i) Write down the exact time when memory will be allocated to each process. (2 marks)



- ii) Indicate the times when external fragmentation is observed in the system. Justify. (1 mark)

~~External fragmentation is never observed at 10:04~~

①

No external fragmentation observed. Every process that waits is due to space limitation of the main memory and not due to the presence of active partitions.

DEPARTMENT OF COMPUTER SCIENCE & ENGINEERING

I midterm examination winter 2011-12

CSU313-Operating systems

A

Time: 1 Hour

Max. Marks: 20

1. The modern systems prefer interrupt driven I/O for slow I/O devices? Why this approach is beneficial for operating systems?

Ans: The above defined statement is ~~carried out to maximise~~⁽²⁾ the CPU utilisation time. Without interrupt driven I/O; suppose the CPU is executing a process; and it request for an I/O device; then the CPU will wait till the request is executed, & Being idle for this period (in terms of sec as slow I/O devices); the CPU utilization is very low. Hence wastage. But in interrupt driven I/O; when it request for an I/O device; the CPU switches to the next process to utilise that time. Once the I/O request is over; interrupt occurs at CPU; CPU leaves the process it is executing (saves) and execute again process 1.

2. What information about a process needs to be saved, changed or updated when a context switch occurs?

Ans: When context switch occurs (saving the current⁽³⁾ process details and moving to the new process); the PC value of the process; i.e. the next command to be executed; accounting information (how much CPU utilisation has already occurred); process id of the process; process state (running/ waiting, etc); are saved of the last process. The registers are loaded for the new process PCB. ~~at that time to the~~ In case of an interrupt; the PC and IP values are updated in the interrupt driven routine table for the corresponding interrupt vector. Now the program counter for changes for the new process to be executed.

3. int main()

```
{  
    int n,m;  
    pid_t pid1,pid2,pid3;  
    n=11;  
    m=12;  
    pid1=fork();  
    if(pid1==0)  
        sh1();
```

```
    execv("edittext.exe",NULL);  
    n=n+15;  
    m=m+10;
```



Q) printf("In first if %d,%d",n,m)

```
{  
    m=m+50;   m = 61  
    n=n+30;   n = 61  
    pid2=fork();  
    if(pid2!=0)  
    {  
        m=m+5; m = 62  
        n=n+5; n = 62  
        pid2=wait(NULL);  
        printf("Second printf %d,%d",n,m);  
    }  
    printf(" Last printf %d,%d",n,m)  
    exit(0)  
}
```

In the above program

- a) How many times the first printf() statement is executed and what are the values of n, m ? Justify.

Ans: First printf() statement is not executed even once.
Executed zero times. As the exec() command for the child implies that child will execute some other code which is edittext.txt; Hence will not even execute the following code :
 $n = n + 15;$
 $m = m + 10;$

printf ("In first if %d,%d", n, m)
So, No values of m, n will be printed. Values actually are still the declared ones.

- b) How many times the second printf() statement is executed and what are the values of n, m ? Justify.

Ans: It will be executed only once by parent as the first child is already executing some other code ; and only left is parent for which pid2 != 0 (given condition); The printed values are : Second printf 46 67 (by parent)
But the parent will wait for any of the two children to finish execution due to pid2 = wait (NULL); before printing .

- c) How many times the last printf() statement is executed and what are the values of n, m ? Justify.

Ans: Two times ; one time by parent and one time by the second child.

When parent execute : → last printf 46
When child 2 execute : → last printf 41
When parent : → after the second printf statement, third printf statement will be executed and hence values will be same ; and child (2) will be executed first time the third printf statement is the first before forking C2 ; and on just before executing any m = 62 , n = 41 ; so those values will be printed .